

# MADRID 2011

```
In [ ]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, LogisticRegression, Lasso, Ridge
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv("C:/Users/user/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_data.csv")
df
```

```
Out[2]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2011-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	84.0	NaN	NaN	NaN	6.0	NaN	NaN	2
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.7	2
2	2011-11-01 01:00:00	2.9	NaN	3.8	NaN	96.0	99.0	NaN	NaN	NaN	NaN	NaN	7.2	2
3	2011-11-01 01:00:00	NaN	0.6	NaN	NaN	60.0	83.0	2.0	NaN	NaN	NaN	NaN	NaN	2
4	2011-11-01 01:00:00	NaN	NaN	NaN	NaN	44.0	62.0	3.0	NaN	NaN	3.0	NaN	NaN	2
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
209923	2011-09-01 00:00:00	NaN	0.2	NaN	NaN	5.0	19.0	44.0	NaN	NaN	NaN	NaN	NaN	2
209924	2011-09-01 00:00:00	NaN	0.1	NaN	NaN	6.0	29.0	NaN	11.0	NaN	7.0	NaN	NaN	2
209925	2011-09-01 00:00:00	NaN	NaN	NaN	0.23	1.0	21.0	28.0	NaN	NaN	NaN	1.44	NaN	2
209926	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	15.0	48.0	NaN	NaN	NaN	NaN	NaN	2
209927	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	4.0	33.0	38.0	13.0	NaN	NaN	NaN	NaN	2

209928 rows × 14 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209928 entries, 0 to 209927
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   date        209928 non-null  object
1   BEN         51393 non-null   float64
2   CO          87127 non-null   float64
3   EBE         51350 non-null   float64
4   NMHC        43517 non-null   float64
5   NO          208954 non-null   float64
6   NO_2        208973 non-null   float64
7   O_3         122049 non-null   float64
8   PM10        103743 non-null   float64
9   PM25        51079 non-null   float64
10  SO_2        87131 non-null   float64
11  TCH         43519 non-null   float64
12  TOL         51175 non-null   float64
13  station     209928 non-null   int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: df1=df.dropna()
df1
```

```
Out[4]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	ε
<b>1</b>	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.7	280
<b>6</b>	2011-11-01 01:00:00	0.7	0.3	1.1	0.16	17.0	66.0	7.0	22.0	16.0	2.0	1.36	1.7	280
<b>25</b>	2011-11-01 02:00:00	1.8	0.3	2.8	0.20	34.0	76.0	3.0	34.0	21.0	8.0	1.71	7.4	280
<b>30</b>	2011-11-01 02:00:00	1.0	0.4	1.3	0.18	31.0	67.0	5.0	25.0	18.0	3.0	1.40	2.9	280
<b>49</b>	2011-11-01 03:00:00	1.3	0.2	2.4	0.22	29.0	72.0	3.0	33.0	20.0	8.0	1.75	6.2	280
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>209862</b>	2011-08-31 22:00:00	0.4	0.1	1.0	0.06	1.0	13.0	33.0	21.0	6.0	5.0	1.26	0.7	280
<b>209881</b>	2011-08-31 23:00:00	0.9	0.1	1.8	0.16	11.0	45.0	30.0	32.0	17.0	3.0	1.34	4.9	280
<b>209886</b>	2011-08-31 23:00:00	0.6	0.1	1.1	0.05	1.0	12.0	48.0	19.0	7.0	5.0	1.26	0.9	280
<b>209905</b>	2011-09-01 00:00:00	0.6	0.1	1.3	0.15	6.0	35.0	34.0	21.0	12.0	3.0	1.32	3.8	280
<b>209910</b>	2011-09-01 00:00:00	0.7	0.1	1.1	0.04	1.0	12.0	46.0	8.0	5.0	5.0	1.25	0.9	280


16460 rows × 14 columns



```
In [5]: df1=df1.drop(["date"],axis=1)
```


```
In [6]: sns.heatmap(df1.corr())
```

```
Out[6]: <AxesSubplot:>
```



```
In [7]: plt.plot(df1["EBE"],df1["PM25"],"o")
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x1aa70adb160>]
```



```
In [8]: x=df1.drop(["EBE"],axis=1)
y=df1["EBE"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


## Linear

```
In [9]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[9]: LinearRegression()
```

```
In [10]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[10]: <matplotlib.collections.PathCollection at 0x1aa70cb4cd0>



```
In [11]: lis=li.score(x_test,y_test)
```

```
In [12]: df1["TCH"].value_counts()
```

Out[12]:

1.30	897
1.29	878
1.28	856
1.31	827
1.27	820
...	
3.41	1
2.88	1
2.41	1
2.80	1
2.49	1

Name: TCH, Length: 171, dtype: int64

```
In [13]: df1.loc[df1["TCH"]<1.40,"TCH"]=1
df1.loc[df1["TCH"]>1.40,"TCH"]=2
df1["TCH"].value_counts()
```

Out[13]:

1.0	12828
2.0	3632

Name: TCH, dtype: int64


## Lasso

```
In [14]: la=Lasso(alpha=5)
la.fit(x_train,y_train)
```

Out[14]: Lasso(alpha=5)

```
In [15]: prediction1=la.predict(x_test)
plt.scatter(y_test,prediction1)
```

Out[15]: <matplotlib.collections.PathCollection at 0x1aa70d1ee80>



```
In [16]: las=la.score(x_test,y_test)
```


## Ridge

```
In [17]: rr=Ridge(alpha=1)
rr.fit(x_train,y_train)
```

Out[17]: Ridge(alpha=1)

```
In [18]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

Out[18]: <matplotlib.collections.PathCollection at 0x1aa70b271f0>



```
In [19]: rrs=rr.score(x_test,y_test)
```


## ElasticNet

```
In [20]: en=ElasticNet()
en.fit(x_train,y_train)
```

Out[20]: ElasticNet()

```
In [21]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

Out[21]: <matplotlib.collections.PathCollection at 0x1aa70d9ca30>



```
In [22]: ens=en.score(x_test,y_test)
```

```
In [23]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

0.8305645679872002

Out[23]: 0.8132155812013233

## Logistic

```
In [24]: g={"TCH":{1.0:"Low",2.0:"High"}}
df1=df1.replace(g)
df1["TCH"].value_counts()
```

Out[24]: Low 12828  
High 3632  
Name: TCH, dtype: int64

```
In [25]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [26]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

Out[26]: LogisticRegression()

```
In [27]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

Out[27]: <matplotlib.collections.PathCollection at 0x1aa70b57cd0>



```
In [28]: los=lo.score(x_test,y_test)
```

## Random Forest

```
In [29]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [30]: g1={"TCH":{"Low":1.0,"High":2.0}}
df1=df1.replace(g1)
```

```
In [31]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [32]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[32]: RandomForestClassifier()



```
In [33]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [34]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[34]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [1, 2, 4, 5, 6],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
    scoring='accuracy')
```

```
In [35]: rfcs=grid_search.best_score_
```

```
In [36]: rfc_best=grid_search.best_estimator_
```

```
In [37]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'])
```

```
Out[37]: [Text(2316.467889908257, 2019.0857142857144, 'O_3 <= 19.5\ngini = 0.338\nsamples = 7259\nvalue = [9040, 2482]\nclass = Yes'),
    Text(1116.0, 1708.457142857143, 'PM25 <= 16.5\ngini = 0.494\nsamples = 1972\nvalue = [1394, 1745]\nclass = No'),
    Text(593.8348623853211, 1397.8285714285716, 'NMHC <= 0.235\ngini = 0.446\nsamples = 1040\nvalue = [1092, 553]\nclass = Yes'),
    Text(327.6330275229358, 1087.2, 'NMHC <= 0.155\ngini = 0.399\nsamples = 915\nvalue = [1043, 396]\nclass = Yes'),
    Text(163.8165137614679, 776.5714285714287, 'NO <= 5.5\ngini = 0.245\nsamples = 219\nvalue = [282, 47]\nclass = Yes'),
    Text(81.90825688073394, 465.9428571428573, 'O_3 <= 12.5\ngini = 0.044\nsamples = 91\nvalue = [130, 3]\nclass = Yes'),
    Text(40.95412844036697, 155.3142857142857, 'gini = 0.117\nsamples = 35\nvalue = [45, 3]\nclass = Yes'),
    Text(122.86238532110092, 155.3142857142857, 'gini = 0.0\nsamples = 56\nvalue = [85, 0]\nclass = Yes'),
    Text(245.72477064220183, 465.9428571428573, 'NO_2 <= 23.5\ngini = 0.348\nsamples = 128\nvalue = [152, 44]\nclass = Yes'),
    Text(204.77064220183485, 155.3142857142857, 'gini = 0.408\nsamples = 12\nvalue = [12, 0]\nclass = Yes')]
```

```
In [38]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

Linear: 0.8305073316874275  
Lasso: 0.5846802065548421  
Ridge: 0.8305645679872002  
ElasticNet: 0.7133987542437089  
Logistic: 0.7778452814904819  
Random Forest: 0.8899496615170978

## Best Model is Random Forest

## MADRID 2012

```
In [39]: df2=pd.read_csv("C:/Users/user/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_data_2011_12.csv")
df2
```

Out[39]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	7.0	18.0	NaN	NaN	NaN	2.0	NaN	NaN	28
1	2012-09-01 01:00:00	0.3	0.3	0.7	NaN	3.0	18.0	55.0	10.0	9.0	1.0	NaN	2.4	28
2	2012-09-01 01:00:00	0.4	NaN	0.7	NaN	2.0	10.0	NaN	NaN	NaN	NaN	NaN	1.5	28
3	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	50.0	NaN	NaN	NaN	NaN	NaN	28
4	2012-09-01 01:00:00	NaN	NaN	NaN	NaN	1.0	13.0	54.0	NaN	NaN	3.0	NaN	NaN	28
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
210715	2012-03-01 00:00:00	NaN	0.6	NaN	NaN	37.0	84.0	14.0	NaN	NaN	NaN	NaN	NaN	28
210716	2012-03-01 00:00:00	NaN	0.4	NaN	NaN	5.0	76.0	NaN	17.0	NaN	7.0	NaN	NaN	28
210717	2012-03-01 00:00:00	NaN	NaN	NaN	0.34	3.0	41.0	24.0	NaN	NaN	NaN	1.34	NaN	28
210718	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	44.0	36.0	NaN	NaN	NaN	NaN	NaN	28
210719	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	56.0	40.0	18.0	NaN	NaN	NaN	NaN	28

210720 rows × 14 columns

```
In [40]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210720 entries, 0 to 210719
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        210720 non-null object
1   BEN         51511 non-null  float64
2   CO          87097 non-null  float64
3   EBE         51482 non-null  float64
4   NMHC        30736 non-null  float64
5   NO          209871 non-null float64
6   NO_2        209872 non-null float64
7   O_3         122339 non-null float64
8   PM10        104838 non-null float64
9   PM25        52164 non-null  float64
10  SO_2        87333 non-null  float64
11  TCH         30736 non-null  float64
12  TOL         51373 non-null  float64
13  station     210720 non-null int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.5+ MB
```

```
In [41]: df3=df2.dropna()  
df3
```

```
Out[41]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	ε
<b>6</b>	2012-09-01 01:00:00	0.4	0.2	0.8	0.24	1.0	7.0	57.0	11.0	7.0	2.0	1.33	0.6	280
<b>30</b>	2012-09-01 02:00:00	0.4	0.2	0.7	0.24	1.0	5.0	55.0	5.0	5.0	2.0	1.33	0.5	280
<b>54</b>	2012-09-01 03:00:00	0.4	0.2	0.7	0.24	1.0	4.0	56.0	6.0	4.0	2.0	1.33	0.5	280
<b>78</b>	2012-09-01 04:00:00	0.3	0.2	0.7	0.25	1.0	5.0	54.0	6.0	5.0	2.0	1.34	0.4	280
<b>102</b>	2012-09-01 05:00:00	0.4	0.2	0.7	0.24	1.0	3.0	53.0	8.0	5.0	2.0	1.33	0.5	280
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>210654</b>	2012-02-29 22:00:00	0.6	0.3	0.5	0.09	1.0	35.0	57.0	25.0	21.0	3.0	1.12	2.3	280
<b>210673</b>	2012-02-29 23:00:00	2.0	0.4	2.4	0.21	16.0	79.0	20.0	37.0	25.0	12.0	1.33	6.2	280
<b>210678</b>	2012-02-29 23:00:00	0.7	0.3	0.6	0.09	1.0	27.0	63.0	22.0	18.0	3.0	1.11	1.9	280
<b>210697</b>	2012-03-01 00:00:00	1.5	0.4	1.7	0.21	16.0	79.0	17.0	28.0	21.0	11.0	1.34	4.9	280
<b>210702</b>	2012-03-01 00:00:00	0.6	0.3	0.5	0.09	1.0	23.0	61.0	18.0	16.0	3.0	1.11	1.2	280


10916 rows × 14 columns



```
In [42]: df3=df3.drop(["date"],axis=1)
```

```
In [43]: sns.heatmap(df3.corr())
```

```
Out[43]: <AxesSubplot:>
```



```
In [44]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear


```
In [45]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[45]: LinearRegression()
```

```
In [ ]:
```

```
In [46]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[46]: <matplotlib.collections.PathCollection at 0x1aa72e176a0>
```



```
In [47]: lis=li.score(x_test,y_test)
```

```
In [48]: df3["TCH"].value_counts()
```

```
Out[48]: 1.30    737
         1.31    676
         1.32    644
         1.33    552
         1.29    529
         ...
         3.03     1
         3.01     1
         2.47     1
         2.33     1
         2.07     1
         Name: TCH, Length: 167, dtype: int64
```

```
In [49]: df3.loc[df3["TCH"]<1.40,"TCH"]=1
         df3.loc[df3["TCH"]>1.40,"TCH"]=2
         df3["TCH"].value_counts()
```

```
Out[49]: 1.0    8772
         2.0    2144
         Name: TCH, dtype: int64
```

```
In [ ]:
```


## Lasso

```
In [50]: la=Lasso(alpha=5)
         la.fit(x_train,y_train)
```

```
Out[50]: Lasso(alpha=5)
```

```
In [51]: prediction1=la.predict(x_test)
plt.scatter(y_test,prediction1)
```

Out[51]: <matplotlib.collections.PathCollection at 0x1aa71741790>



```
In [52]: las=la.score(x_test,y_test)
```


## Ridge

```
In [53]: rr=Ridge(alpha=1)
rr.fit(x_train,y_train)
```

Out[53]: Ridge(alpha=1)

```
In [54]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

Out[54]: <matplotlib.collections.PathCollection at 0x1aa717a3310>



```
In [55]: rrs=rr.score(x_test,y_test)
```




## ElasticNet

```
In [56]: en=ElasticNet()
en.fit(x_train,y_train)
```

Out[56]: ElasticNet()

```
In [57]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

Out[57]: <matplotlib.collections.PathCollection at 0x1aa717f88e0>



```
In [58]: ens=en.score(x_test,y_test)
```

```
In [59]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

0.6854937601485451

Out[59]: 0.6889077427743225

## Logistic

```
In [60]: g={"TCH":{1.0:"Low",2.0:"High"}}
df3=df3.replace(g)
df3["TCH"].value_counts()
```

Out[60]: Low 8772  
High 2144  
Name: TCH, dtype: int64

```
In [61]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [62]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

Out[62]: LogisticRegression()

```
In [63]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

Out[63]: <matplotlib.collections.PathCollection at 0x1aa716bc580>



```
In [64]: los=lo.score(x_test,y_test)
```

## Random Forest

```
In [65]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [66]: g1={"TCH":{"Low":1.0,"High":2.0}}
df3=df3.replace(g1)
```

```
In [67]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [68]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[68]: RandomForestClassifier()

```
In [69]: parameter={
            'max_depth': [1, 2, 4, 5, 6],
            'min_samples_leaf': [5, 10, 15, 20, 25],
            'n_estimators': [10, 20, 30, 40, 50]
        }
```

```
In [70]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accu
grid_search.fit(x_train,y_train)
```

```
Out[70]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 4, 5, 6],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [71]: rfcs=grid_search.best_score_
```

```
In [72]: rfc_best=grid_search.best_estimator_
```

```
In [73]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'])
```

```
Out[73]: [Text(2493.5625, 2019.0857142857144, 'TOL <= 7.15\ngini = 0.308\nsamples = 4812\nvalue = [6188, 1453]\nclass = Yes'),
Text(1488.0, 1708.457142857143, 'NO <= 22.5\ngini = 0.259\nsamples = 4518\nvalue = [6094, 1098]\nclass = Yes'),
Text(744.0, 1397.8285714285716, 'O_3 <= 23.5\ngini = 0.178\nsamples = 3782\nvalue = [5456, 597]\nclass = Yes'),
Text(372.0, 1087.2, 'NMHC <= 0.275\ngini = 0.452\nsamples = 574\nvalue = [603, 317]\nclass = Yes'),
Text(186.0, 776.5714285714287, 'PM10 <= 16.5\ngini = 0.258\nsamples = 406\nvalue = [556, 100]\nclass = Yes'),
Text(93.0, 465.9428571428573, 'station <= 28079016.0\ngini = 0.18\nsamples = 248\nvalue = [361, 40]\nclass = Yes'),
Text(46.5, 155.3142857142857, 'gini = 0.0\nsamples = 44\nvalue = [73, 0]\nclass = Yes'),
Text(139.5, 155.3142857142857, 'gini = 0.214\nsamples = 204\nvalue = [288, 40]\nclass = Yes'),
Text(279.0, 465.9428571428573, 'TOL <= 1.05\ngini = 0.36\nsamples = 158\nvalue = [195, 60]\nclass = Yes'),
Text(232.5, 155.3142857142857, 'gini = 0.5\nsamples = 23\nvalue = [20, 20]
```

```
In [74]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.6851032563985955
Lasso: -0.0014194828338582877
Ridge: 0.6854937601485451
ElasticNet: 0.34195759108298374
Logistic: 0.8006106870229007
Random Forest: 0.9331240896615699
```

## Best model is Random Forest

```
In [ ]:
```