

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, LogisticRegression, Lasso, Ridge, ElasticNet
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv("/Users/bhoomish/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_2013.csv")
df
```

```
Out[2]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2013-11-01 01:00:00	NaN	0.6	NaN	NaN	135.0	74.0	NaN	NaN	NaN	7.0	NaN	NaN	28079004
1	2013-11-01 01:00:00	1.5	0.5	1.3	NaN	71.0	83.0	2.0	23.0	16.0	12.0	NaN	8.3	28079008
2	2013-11-01 01:00:00	3.9	NaN	2.8	NaN	49.0	70.0	NaN	NaN	NaN	NaN	NaN	9.0	28079011
3	2013-11-01 01:00:00	NaN	0.5	NaN	NaN	82.0	87.0	3.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2013-11-01 01:00:00	NaN	NaN	NaN	NaN	242.0	111.0	2.0	NaN	NaN	12.0	NaN	NaN	28079017
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
209875	2013-03-01 00:00:00	NaN	0.4	NaN	NaN	8.0	39.0	52.0	NaN	NaN	NaN	NaN	NaN	28079056
209876	2013-03-01 00:00:00	NaN	0.4	NaN	NaN	1.0	11.0	NaN	6.0	NaN	2.0	NaN	NaN	28079057
209877	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	4.0	75.0	NaN	NaN	NaN	NaN	NaN	28079058
209878	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	11.0	52.0	NaN	NaN	NaN	NaN	NaN	28079059
209879	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	1.0	10.0	75.0	3.0	NaN	NaN	NaN	NaN	28079060

209880 rows × 14 columns

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209880 entries, 0 to 209879
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209880 non-null  object
1   BEN         50462 non-null   float64
2   CO          87018 non-null   float64
3   EBE         50463 non-null   float64
4   NMHC        25935 non-null   float64
5   NO          209108 non-null  float64
6   NO_2        209108 non-null  float64
7   O_3         121858 non-null  float64
8   PM10        104339 non-null  float64
9   PM25        51980 non-null   float64
10  SO_2        86970 non-null   float64
11  TCH         25935 non-null   float64
12  TOL         50317 non-null   float64
13  station     209880 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: df1=df.dropna()  
df1
```

```
Out[4]:
```


	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
17286	2013-08-01 01:00:00	0.4	0.2	0.8	0.28	1.0	24.0	79.0	35.0	8.0	3.0	1.49	1.3	28079024
17310	2013-08-01 02:00:00	0.5	0.2	0.9	0.28	1.0	16.0	93.0	60.0	18.0	3.0	1.61	4.0	28079024
17334	2013-08-01 03:00:00	0.5	0.2	1.1	0.29	1.0	14.0	90.0	38.0	12.0	3.0	1.71	2.8	28079024
17358	2013-08-01 04:00:00	0.6	0.2	1.2	0.26	1.0	12.0	84.0	30.0	8.0	3.0	1.44	2.8	28079024
17382	2013-08-01 05:00:00	0.3	0.2	0.8	0.25	1.0	15.0	72.0	25.0	7.0	3.0	1.40	1.7	28079024
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
209622	2013-02-28 14:00:00	1.1	0.3	0.3	0.27	3.0	17.0	64.0	5.0	5.0	2.0	1.41	0.9	28079024
209646	2013-02-28 15:00:00	1.3	0.4	0.3	0.27	2.0	16.0	66.0	6.0	5.0	1.0	1.40	0.9	28079024
209670	2013-02-28 16:00:00	1.1	0.3	0.3	0.27	1.0	17.0	65.0	5.0	4.0	1.0	1.40	0.7	28079024
209694	2013-02-28 17:00:00	1.0	0.3	0.4	0.27	1.0	18.0	64.0	5.0	5.0	1.0	1.39	0.7	28079024
209718	2013-02-28 18:00:00	1.0	0.3	0.4	0.27	1.0	22.0	62.0	6.0	6.0	1.0	1.39	0.7	28079024

7315 rows × 14 columns

```
In [5]: df1=df1.drop(["date"],axis=1)
```


```
sns.heatmap(df1.corr())
```

```
Out[6]: <Axes: >
```



```
In [7]: plt.plot(df1["EBE"],df1["NMHC"],"o")
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x7f9d822f5150>]
```



```
In [8]: data=df[["EBE","NMHC"]]
```

```
In [9]: x=df1.drop(["EBE"],axis=1)  
y=df1["EBE"]  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear

```
In [10]: li=LinearRegression()  
li.fit(x_train,y_train)
```


Out[10]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [11]: prediction=li.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[11]: <matplotlib.collections.PathCollection at 0x7f9d92470940>



```
In [12]: lis=li.score(x_test,y_test)
```

```
In [13]: df1["TCH"].value_counts()
```

```
Out[13]: 1.32    888
         1.33    843
         1.34    729
         1.31    719
         1.35    556
         ...
         2.39     1
         2.22     1
         2.29     1
         2.38     1
         2.80     1
         Name: TCH, Length: 114, dtype: int64
```

```
In [14]: df1.loc[df1["TCH"]<1.40, "TCH"]=1
         df1.loc[df1["TCH"]>1.40, "TCH"]=2
         df1["TCH"].value_counts()
```

```
Out[14]: 1.0    5718
         2.0    1597
         Name: TCH, dtype: int64
```

## Lasso


```
In [15]: la=Lasso(alpha=5)
         la.fit(x_train,y_train)
```

```
Out[15]: Lasso(alpha=5)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [16]: prediction1=la.predict(x_test)
plt.scatter(y_test,prediction1)
```

```
Out[16]: <matplotlib.collections.PathCollection at 0x7f9d924dab30>
```



```
In [17]: las=la.score(x_test,y_test)
```

## Ridge



```
In [18]: rr=Ridge(alpha=1)
rr.fit(x_train,y_train)
```


Out[18]: Ridge(alpha=1)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [19]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

Out[19]: <matplotlib.collections.PathCollection at 0x7f9d9256b550>



```
In [20]: rrs=rr.score(x_test,y_test)
```

## ElasticNet


```
In [21]: en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[21]: ElasticNet()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [22]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[22]: <matplotlib.collections.PathCollection at 0x7f9d9240b4f0>
```



```
In [23]: ens=en.score(x_test,y_test)
```

```
In [24]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.3923674796819595
```

```
Out[24]: 0.3927671300897322
```

# Logistic

```
In [25]: g={"TCH":{1.0:"Low",2.0:"High"}}  
df1=df1.replace(g)  
df1["TCH"].value_counts()
```

```
Out[25]: Low      5718  
        High     1597  
        Name: TCH, dtype: int64
```

```
In [26]: x=df1.drop(["TCH"],axis=1)  
        y=df1["TCH"]  
        x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


```
In [27]: lo=LogisticRegression()  
        lo.fit(x_train,y_train)
```

```
Out[27]: LogisticRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [28]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[28]: <matplotlib.collections.PathCollection at 0x7f9d30433c70>
```



```
In [29]: los=lo.score(x_test,y_test)
```

## Random Forest

```
In [30]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [31]: g1={"TCH":{"Low":1.0,"High":2.0}}
df1=df1.replace(g1)
```

```
In [32]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [33]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[33]: RandomForestClassifier()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [34]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [35]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[35]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 4, 5, 6],  
'min\_samples\_leaf': [5, 10, 15, 20, 25],  
'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [36]: rfcs=grid_search.best_score_
```

```
In [37]: rfc_best=grid_search.best_estimator_
```

```
In [38]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[38]: [Text(0.5375, 0.9285714285714286, 'NMHC <= 0.275\ngini = 0.328\nsamples = 3238\nvalue = [4060, 1060]\ncla  
ss = Yes'),  
Text(0.3390625, 0.7857142857142857, 'TOL <= 1.05\ngini = 0.09\nsamples = 2599\nvalue = [3927, 194]\nclas  
s = Yes'),  
Text(0.196875, 0.6428571428571429, 'O_3 <= 45.5\ngini = 0.031\nsamples = 1960\nvalue = [3090, 50]\nnclass  
= Yes'),  
Text(0.1, 0.5, 'PM10 <= 17.5\ngini = 0.174\nsamples = 215\nvalue = [309, 33]\nnclass = Yes'),  
Text(0.05, 0.35714285714285715, 'O_3 <= 20.5\ngini = 0.123\nsamples = 191\nvalue = [283, 20]\nnclass = Ye  
s'),  
Text(0.025, 0.21428571428571427, 'SO_2 <= 1.5\ngini = 0.335\nsamples = 30\nvalue = [37, 10]\nnclass = Ye  
s'),  
Text(0.0125, 0.07142857142857142, 'gini = 0.497\nsamples = 9\nvalue = [7, 6]\nnclass = Yes'),  
Text(0.0375, 0.07142857142857142, 'gini = 0.208\nsamples = 21\nvalue = [30, 4]\nnclass = Yes'),  
Text(0.075, 0.21428571428571427, 'PM10 <= 5.5\ngini = 0.075\nsamples = 161\nvalue = [246, 10]\nnclass = Y  
es'),  
Text(0.0625, 0.07142857142857142, 'gini = 0.0\nsamples = 46\nvalue = [75, 0]\nnclass = Yes'),  
Text(0.0875, 0.07142857142857142, 'gini = 0.104\nsamples = 115\nvalue = [171, 10]\nnclass = Yes'),  
Text(0.15, 0.35714285714285715, 'O_3 <= 37.5\ngini = 0.444\nsamples = 24\nvalue = [26, 13]\nnclass = Ye  
s'),  
Text(0.125, 0.31428571428571427, 'PM25 <= 10.5\ngini = 0.408\nsamples = 123\nvalue = [50, 61]\nnclass = Y
```

```
In [39]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.40767907025541705
Lasso: -0.00029951584872067727
Ridge: 0.3923674796819595
ElasticNet: 0.10978052070818145
Logistic: 0.7890660592255125
Random Forest: 0.946484375
```

**Best Model is Random Forest**



```
In [40]: df2=pd.read_csv("/Users/bob/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_2014.csv")
df2
```

Out[40]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2014-06-01 01:00:00	NaN	0.2	NaN	NaN	3.0	10.0	NaN	NaN	NaN	3.0	NaN	NaN	28079004
1	2014-06-01 01:00:00	0.2	0.2	0.1	0.11	3.0	17.0	68.0	10.0	5.0	5.0	1.36	1.3	28079008
2	2014-06-01 01:00:00	0.3	NaN	0.1	NaN	2.0	6.0	NaN	NaN	NaN	NaN	NaN	1.1	28079011
3	2014-06-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	79.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2014-06-01 01:00:00	NaN	NaN	NaN	NaN	1.0	6.0	75.0	NaN	NaN	4.0	NaN	NaN	28079017
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
210019	2014-09-01 00:00:00	NaN	0.5	NaN	NaN	20.0	84.0	29.0	NaN	NaN	NaN	NaN	NaN	28079056
210020	2014-09-01 00:00:00	NaN	0.3	NaN	NaN	1.0	22.0	NaN	15.0	NaN	6.0	NaN	NaN	28079057
210021	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	1.0	13.0	70.0	NaN	NaN	NaN	NaN	NaN	28079058
210022	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	38.0	42.0	NaN	NaN	NaN	NaN	NaN	28079059
210023	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	1.0	26.0	65.0	11.0	NaN	NaN	NaN	NaN	28079060

210024 rows × 14 columns

```
In [41]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210024 entries, 0 to 210023
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        210024 non-null  object
1   BEN         46703 non-null   float64
2   CO          87023 non-null   float64
3   EBE         46722 non-null   float64
4   NMHC        25021 non-null   float64
5   NO          209154 non-null  float64
6   NO_2        209154 non-null  float64
7   O_3         121681 non-null  float64
8   PM10        104311 non-null  float64
9   PM25        51954 non-null   float64
10  SO_2        87141 non-null   float64
11  TCH         25021 non-null   float64
12  TOL         46570 non-null   float64
13  station     210024 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [42]: df3=df2.dropna()  
df3
```

```
Out[42]:
```


	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
1	2014-06-01 01:00:00	0.2	0.2	0.1	0.11	3.0	17.0	68.0	10.0	5.0	5.0	1.36	1.3	28079008
6	2014-06-01 01:00:00	0.1	0.2	0.1	0.23	1.0	5.0	80.0	4.0	3.0	2.0	1.21	0.1	28079024
25	2014-06-01 02:00:00	0.2	0.2	0.1	0.11	4.0	21.0	63.0	9.0	6.0	5.0	1.36	0.8	28079008
30	2014-06-01 02:00:00	0.2	0.2	0.1	0.23	1.0	4.0	88.0	7.0	5.0	2.0	1.21	0.1	28079024
49	2014-06-01 03:00:00	0.1	0.2	0.1	0.11	4.0	18.0	66.0	9.0	7.0	6.0	1.36	0.9	28079008
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
209958	2014-08-31 22:00:00	0.2	0.2	0.1	0.22	1.0	28.0	96.0	61.0	15.0	3.0	1.28	0.1	28079024
209977	2014-08-31 23:00:00	1.1	0.7	0.7	0.19	36.0	118.0	23.0	60.0	25.0	9.0	1.27	6.5	28079008
209982	2014-08-31 23:00:00	0.2	0.2	0.1	0.21	1.0	17.0	90.0	28.0	14.0	3.0	1.27	0.2	28079024
210001	2014-09-01 00:00:00	0.6	0.4	0.4	0.12	6.0	63.0	41.0	26.0	15.0	8.0	1.19	4.1	28079008
210006	2014-09-01 00:00:00	0.2	0.2	0.1	0.23	1.0	30.0	69.0	18.0	13.0	3.0	1.30	0.1	28079024

13946 rows × 14 columns

```
In [43]: df3=df3.drop(["date"],axis=1)
```

```
In [44]: sns.heatmap(df3.corr())
```

```
Out[44]: <Axes: >
```



```
In [45]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear


```
In [46]: li=LinearRegression()  
li.fit(x_train,y_train)
```

Out[46]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [47]: prediction=li.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[47]: <matplotlib.collections.PathCollection at 0x7f9d82aff730>



```
In [48]: lis=li.score(x_test,y_test)
```

```
In [49]: df3["TCH"].value_counts()
```

```
Out[49]: 1.37    601
         1.36    598
         1.34    529
         1.35    528
         1.38    515
         ...
         4.39     1
         4.08     1
         3.42     1
         2.98     1
         2.69     1
         Name: TCH, Length: 184, dtype: int64
```

```
In [50]: df3.loc[df3["TCH"]<1.40,"TCH"]=1
         df3.loc[df3["TCH"]>1.40,"TCH"]=2
         df3["TCH"].value_counts()
```

```
Out[50]: 1.0    9997
         2.0    3949
         Name: TCH, dtype: int64
```

```
In [ ]:
```

## Lasso


```
In [51]: la=Lasso(alpha=5)
         la.fit(x_train,y_train)
```

```
Out[51]: Lasso(alpha=5)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [52]: prediction1=la.predict(x_test)
plt.scatter(y_test,prediction1)
```

```
Out[52]: <matplotlib.collections.PathCollection at 0x7f9d8297ceb0>
```



```
In [53]: las=la.score(x_test,y_test)
```

## Ridge

```
In [54]: rr=Ridge(alpha=1)
rr.fit(x_train,y_train)
```


Out[54]: Ridge(alpha=1)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [55]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

Out[55]: <matplotlib.collections.PathCollection at 0x7f9d829cc9d0>





```
In [56]: rrs=rr.score(x_test,y_test)
```

## ElasticNet


```
In [57]: en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[57]: ElasticNet()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [58]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[58]: <matplotlib.collections.PathCollection at 0x7f9d82b49030>
```



```
In [59]: ens=en.score(x_test,y_test)
```

```
In [60]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.7091386251261051
```

```
Out[60]: 0.7047426795616607
```

# Logistic

```
In [61]: g={"TCH":{1.0:"Low",2.0:"High"}}  
df3=df3.replace(g)  
df3["TCH"].value_counts()
```

```
Out[61]: Low      9997  
        High     3949  
        Name: TCH, dtype: int64
```

```
In [62]: x=df3.drop(["TCH"],axis=1)  
        y=df3["TCH"]  
        x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [63]: lo=LogisticRegression()  
        lo.fit(x_train,y_train)
```

```
Out[63]: LogisticRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [64]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[64]: <matplotlib.collections.PathCollection at 0x7f9d3046cb80>
```



```
In [65]: los=lo.score(x_test,y_test)
```

## Random Forest

```
In [66]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [67]: g1={"TCH":{"Low":1.0,"High":2.0}}
df3=df3.replace(g1)
```

```
In [68]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [69]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[69]: RandomForestClassifier()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [70]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [71]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[71]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 4, 5, 6],  
'min\_samples\_leaf': [5, 10, 15, 20, 25],  
'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [72]: rfcs=grid_search.best_score_
```

```
In [73]: rfc_best=grid_search.best_estimator_
```

```
In [74]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[74]: [Text(0.5191326530612245, 0.9285714285714286, 'PM25 <= 12.5\ngini = 0.402\nsamples = 6171\nvalue = [7045,  
2717]\nnclass = Yes'),  
Text(0.2755102040816326, 0.7857142857142857, 'NMHC <= 0.285\ngini = 0.289\nsamples = 4464\nvalue = [583  
6, 1237]\nnclass = Yes'),  
Text(0.14540816326530612, 0.6428571428571429, 'SO_2 <= 8.5\ngini = 0.233\nsamples = 4206\nvalue = [5776,  
898]\nnclass = Yes'),  
Text(0.08163265306122448, 0.5, 'NO_2 <= 25.5\ngini = 0.211\nsamples = 4094\nvalue = [5722, 779]\nnclass =  
Yes'),  
Text(0.04081632653061224, 0.35714285714285715, 'NO_2 <= 11.5\ngini = 0.084\nsamples = 2607\nvalue = [395  
3, 182]\nnclass = Yes'),  
Text(0.02040816326530612, 0.21428571428571427, 'NO_2 <= 4.5\ngini = 0.023\nsamples = 1603\nvalue = [253  
6, 30]\nnclass = Yes'),  
Text(0.01020408163265306, 0.07142857142857142, 'gini = 0.008\nsamples = 800\nvalue = [1275, 5]\nnclass =  
Yes'),  
Text(0.030612244897959183, 0.07142857142857142, 'gini = 0.038\nsamples = 803\nvalue = [1261, 25]\nnclass  
= Yes'),  
Text(0.061224489795918366, 0.21428571428571427, 'NMHC <= 0.135\ngini = 0.175\nsamples = 1004\nvalue = [1  
417, 152]\nnclass = Yes'),  
Text(0.05102040816326531, 0.07142857142857142, 'gini = 0.05\nsamples = 279\nvalue = [418, 11]\nnclass = Y  
,\n',\n'
```

```
In [75]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.7105902856424091
Lasso: -0.0001259044549823951
Ridge: 0.7091386251261051
ElasticNet: 0.43942661468297217
Logistic: 0.7127151051625239
Random Forest: 0.8929522638803524
```

**Best model is Random Forest**